

# HW1

January 30, 2023

```
[ ]: import autograd.numpy as np
      from autograd.misc.flatten import flatten_func
      from autograd import grad
      import matplotlib.pyplot as plt
```

## 0.1 Problem 1 (13.1)

```
[ ]: data = np.loadtxt('2_eggs.csv', delimiter=',')
      x = data[:2,:]
      y = data[2,:][np.newaxis,:]

      print(np.shape(x))
      print(np.shape(y))
```

(2, 96)

(1, 96)

```
[ ]: def normalize(input):
      for i in range(np.shape(input)[0]):
          input[i] = (input[i] - np.mean(input[i]))/np.std(input[i])
      return input
```

```
[ ]: x = normalize(x)
```

```
[ ]: def feature_transforms(a, w):
      for W in w:

          # compute inner-product with current layer weights
          a = W[0] + np.dot(a.T, W[1:])

          # pass through activation
          a = np.tanh(a).T

      return a
```

```
[ ]: # neural network model
      def model(x, theta):
          # compute feature transformation
```

```

f = feature_transforms(x, theta[0])

# compute final linear combination
a = theta[1][0] + np.dot(f.T, theta[1][1:])

return a.T

```

```

[ ]: # create initial weights for a neural network model
def network_initializer(layer_sizes, scale):

    # container for all tunable weights
    weights = []

    # create appropriately -sized initial weight matrix for each layer of
    ↪ network
    for k in range(len(layer_sizes)-1):

        # get layer sizes for current weight matrix
        U_k = layer_sizes[k]
        U_k_plus_1 = layer_sizes[k+1]

        # make weight matrix
        weight = scale*np.random.randn(U_k+1, U_k_plus_1)
        weights.append(weight)

    # repackage weights so that theta_init[0] contains all
    # weight matrices internal to the network, and theta_init[1]
    # contains final linear combination weights
    theta_init = [weights[:-1], weights[-1]]

    return theta_init

```

```

[ ]: layer_sizes = [2, 10, 10, 10, 10, 1]
theta = network_initializer(layer_sizes, 1)

```

```

[ ]: def softmax(w):
    cost = np.sum(np.log(1 + np.exp(-y*model(x,w))))
    return cost/float(np.size(y))

```

```

[ ]: def gradient_descent(g, alpha, max_its, w):
    weight_history = [w]
    cost_history = [g(w)]

    g_flat, unflatten, w = flatten_func(g, w)
    gradient = grad(g_flat)

    for k in range(max_its):

```

```

grad_eval = gradient(w)

w = w - alpha*grad_eval

weight_history.append(unflatten(w))
cost_history.append(g(unflatten(w)))

return weight_history, cost_history

```

```

[ ]: alpha = 0.3
num_iter = 1000
weight_history, cost_history = gradient_descent(softmax, alpha, num_iter, theta)

```

```


[ ]: plt.plot([np.sum(np.equal(np.sign(model(x, w)), y))/np.size(y) for w in
↪weight_history])
plt.xlabel("Number of iterations")
plt.ylabel("Accuracy")

```

```


[ ]: Text(0, 0.5, 'Accuracy')

```



```
[ ]: plt.plot(cost_history)
plt.xlabel("Number of iterations")
plt.ylabel("Cost")
```

```
[ ]: Text(0, 0.5, 'Cost')
```



We can see that our classification accuracy converges to 1 and our cost function converges to 0

## 0.2 Problem 2 (13.2)

```
[ ]: data = np.loadtxt('3_layercake_data.csv', delimiter=',')
x = data[:2,:]
y = data[2,:][np.newaxis,:]

print(np.shape(x))
print(np.shape(y))
```

```
(2, 110)
```

```
(1, 110)
```

```
[ ]: x = normalize(x)
```

```
[ ]: layer_sizes = [2, 12, 5, 3]
theta = network_initializer(layer_sizes, 1)


[ ]: def multi_softmax(w):

    model_output = model(x,w)
    cost = np.sum(np.log(np.sum(np.exp(model_output), axis=0)) - np.
↪array([model_output[int(y[0][i]), i] for i in range(len(y[0]))]))
    return cost/float(np.size(y))

[ ]: alpha = 0.4
num_iter = 1000
weight_history, cost_history = gradient_descent(multi_softmax, alpha, num_iter,
↪theta)


[ ]: plt.plot([np.sum(np.equal(np.argmax(model(x, w), axis=0), y))/np.size(y) for w
↪in weight_history])
plt.xlabel("Number of iterations")
plt.ylabel("Accuracy")

[ ]: Text(0, 0.5, 'Accuracy')
```



```
[ ]: plt.plot(cost_history)
plt.xlabel("Number of iterations")
plt.ylabel("Cost")
```

```
[ ]: Text(0, 0.5, 'Cost')
```



We can see that our classification accuracy converges to 1 and our cost function converges to 0

### 0.3 Problem 3 (13.3)

#### 0.3.1 (a)

Between any two layers (with number of units  $N$  and  $M$ , respectively) in a general feed-forward neural network, there are  $N \times (M+1)$  tunable parameters corresponding to each connection between units ( $N \times M$ ) and the connections between the bias term of the first layer and each unit in the second layer ( $M$ ). Therefore, for a general  $L$ -hidden-layer neural network where  $N$  is the input dimension,  $U_1$  through  $U_L$  are the number of desired units in the hidden layers 1 through  $L$ , respectively, and  $C$  is the output dimension, we have the total number of tunable parameters

$$Q = N(U_1 + 1) + \sum_{i=1}^{L-1} U_i(U_{i+1} + 1) + U_L(C + 1)$$

### 0.3.2 (b)

The input dimension  $N$  contributes to  $Q$  only through a linear relationship with the number of units in the first hidden layer ( $U_1$ ), while the number of datapoints  $P$  does not contribute to  $Q$  at all.